Adaptive Truncation and Simulation-Based Calibration Applied to COM-Poisson Models

Wellington José Leite da Silva

Rio de Janeiro, RJ
2026

FGV EMAp
ESCOLA DE MATEMÁTICA APLICADA

The Challenge of Count Data

Count data are ubiquitous across various scientific and social domains, representing the frequency of occurrence of discrete events.

- **Common Examples:**
  - Number of website visits.
  - Purchases in physical or online stores.
  - Calls to a call center.
  - Disease cases in a population.
  - Species abundance in a habitat.
- **Real-world data often exhibit:**
  - Equidispersion (Variance = Mean)
  - Overdispersion (Variance > Mean)
  - Underdispersion (Variance < Mean)
- **Challenge:** Accurately modelling the inherent variability in these data.

- Hilbe, 2012
- Sellers, 2021
- Cameron & Trivedi, 2013

The Poisson Distribution: Fundamentals

FGV EMAp ESCOLA DE MATEMÁTICA APLICADA

The Poisson distribution has been the standard choice for modelling count data due to its simplicity and interpretability.

- **Probability Mass Function (PMF):**

$$P(X = k) = \frac{\lambda^k e^{-\lambda}}{k!}$$

  Where $\lambda > 0$ is the rate parameter.

- **Key Assumption: Equidispersion**
  - The mean and variance are equal: $E[X] = Var[X] = \lambda$.
  - This is a defining characteristic of the Poisson distribution.

▶ Shmueli et al., 2005

The Need for Flexible Models

To overcome the limitations of the Poisson distribution, especially its restrictive equidispersion assumption, more flexible models are needed.

- **Alternative Models:** Negative Binomial distribution (for over-dispersion), zero-inflated models, etc.
- **Challenge:** Finding a model that can uniformly accommodate both over-dispersion and under-dispersion.
- **The Path Forward:** The Conway-Maxwell-Poisson (COM-Poisson) distribution emerges as a promising solution.

▶ Hilbe, 2012

The Conway-Maxwell-Poisson (COM-Poisson) Distribution

FGV EMAp ESCOLA DE MATEMÁTICA APLICADA

The COM-Poisson or CMP distribution generalises the Poisson by incorporating a dispersion parameter (ν), allowing it to model a wide range of dispersion patterns.

- **Origin:** Introduced by Conway and Maxwell (1962) in the context of queuing theory.
- **Flexibility:** Capable of capturing:
  - Under-dispersion ($Var[X] < E[X]$)
  - Equidispersion ($Var[X] = E[X]$, reducing to Poisson when $\nu = 1$)
  - Over-dispersion ($Var[X] > E[X]$)
- **Parameters:** $\lambda > 0$ (rate parameter) and $\nu > 0$ (dispersion parameter).
- **Usual reparametrisation:**
  - $\mu := \lambda^{1/\nu}$ (asymptotic approximation for $E[X]$), $\nu > 0$
  - Important for regression modelling

- Shmueli et al., 2005
- Conway & Maxwell, 1962

Dispersion Regimes of the COM-Poisson

The flexibility of the COM-Poisson is evident in its ability to adapt to different dispersion regimes, depending on the values of ν.

Parameters	**Resulting Distribution**	**Dispersion Type**
$\nu = 1$	Poisson	Equidispersion (Variance = Mean)
$\lambda < 1,\ \nu \to 0$	Geometric	Over-dispersion (Variance > Mean)
$\nu \to \infty$	Bernoulli	Under-dispersion (Variance < Mean)
$\nu > 1$	COM-Poisson	Under-dispersion (Variance < Mean)
$\nu < 1$	COM-Poisson	Over-dispersion (Variance > Mean)

Tabela: Dispersion regimes of the COM-Poisson distribution.

Probability Mass Function of the COM-Poisson

The PMF of the COM-Poisson is defined by two parameters, λ and ν, and includes a crucial normalising constant.

- **COM-Poisson PMF:**

$$P(X = k \mid \lambda, \nu) = \frac{\lambda^k}{(k!)^\nu Z(\lambda, \nu)}, \quad k = 0, 1, 2, \ldots$$

where

$$Z(\lambda, \nu) = \sum_{i=0}^{\infty} \frac{\lambda^i}{(i!)^\nu}$$

- **Problem:** $Z(\lambda, \nu)$ has no closed form (except when $\nu = 1$)

▶ Shmueli et al., 2005

Adaptive Truncation

Adaptive truncation strategies provide a principled approach for approximating infinite sums with certified error bounds.

- **Core Idea:** Approximate an infinite sum by truncating it at a certain point, while providing a guarantee on the maximum error introduced by this truncation.
- **Benefit:** Allows for arbitrary precision (ε) in the computation of the normalising constant, crucial for robust statistical modelling.

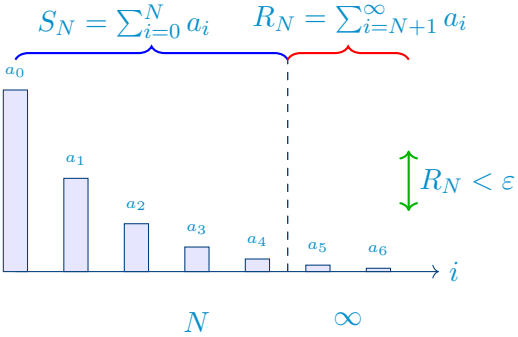

- Braden, 1992
- Carvalho, Silva, and Moreira, 2025

Adaptive Truncation

We propose 3 approaches, for a infinite sum $S = \sum_{n=0}^{\infty} a_n$:

- **Ratio-Bounding Pair**. Based on the ratio test for convergent series, providing explicit error bounds.
- **Integral-Bounding Pair**. Uses integrals to majorize and minorize the remainder of the series, suitable for functions that can be integrated, and has closed form.
- **Sum-to-threshold**. Truncates the series when the remaining terms fall below a certain threshold, only guaranteed to $n > n_0$, here n_0 can be very large.

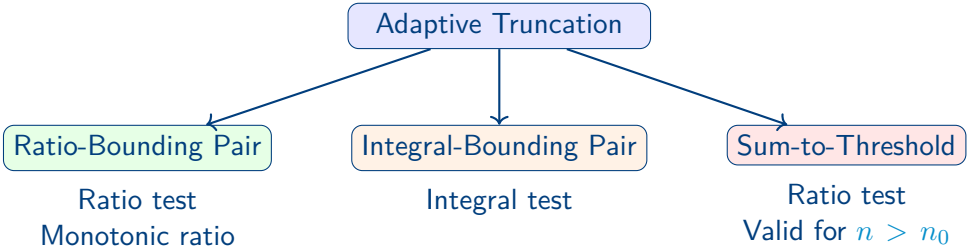

- Braden, 1992
- Carvalho, Silva, and Moreira, 2025

Adaptive Truncation

FGV EMAp
ESCOLA DE MATEMÁTICA APLICADA

Theorem 1: Ratio Bounding Pair dominates Sum-to-threshold

Let $\{a_n\}$ be positive, decreasing, and satisfy the ratio test with limit $L < 1$. If $r_n = \frac{a_{n+1}}{a_n}$ is monotonic, i.e., we can apply both approaches, Ratio Bounding Pair and Sum-to-Threshold. Then, the Ratio Bounding Pair approach uses fewer terms compared to the Sum-to-Threshold approach at the same precision level.

Proof of Theorem 1

See Appendix A.

- Carvalho, Silva, and Moreira, 2025

Ratio Bounding Pair approach

The applicability of the Ratio Bounding Pair depends on the series satisfying the ratio test, which is a common condition for the convergence of many series.

- **Ratio Test Condition:** A positive series $\sum_{n=0}^{\infty} a_n$ converges if the limit of the ratio of consecutive terms is less than 1, i.e. $\lim_{n\to\infty} \frac{a_{n+1}}{a_n} = L < 1$.

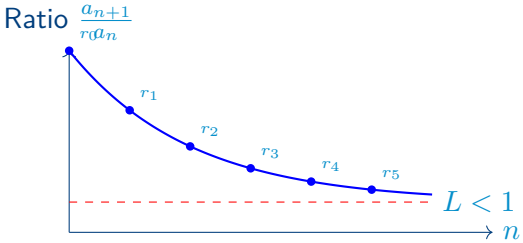

The ratio of consecutive terms must converge to a value $L < 1$.

Ratio Bounding Pair approach

The main result formalises how to construct the error bounds for a series that satisfies the ratio test, providing the tools for adaptive truncation.

- **Context:** Given a series $S = \sum_{k=0}^{\infty} a_k$ with $a_k > 0$ that satisfies the ratio test with limit L.
- **Bounds:** If the ratio a_{n+1}/a_n is monotonically decreasing for L:

$$\frac{a_{n+1}}{1-L} < S - S_n < \frac{a_{n+1}}{1 - a_{n+1}/a_n}$$

  If the ratio a_{n+1}/a_n is monotonically increasing for L:

$$\frac{a_{n+1}}{1 - a_{n+1}/a_n} < S - S_n < \frac{a_{n+1}}{1-L}$$

- **Implication:** These bounds allow estimating the error $S - S_n$ and ensuring it is within a tolerance ε.

▶ Carvalho, Silva, and Moreira, 2025

Ratio Bounding Pair for COM-Poisson

The Ratio Bounding Pair approach is particularly suitable for the COM-Poisson normalising constant due to the structure of its terms.

- **COM-Poisson Series Terms:** The normalising constant $Z(\lambda, \nu) = \sum_{i=0}^{\infty} a_i$ is defined by the sequence $a_n = \frac{\lambda^n}{(n!)^\nu}, \quad n \geq 0.$
- **Ratio of Consecutive Terms:** The ratio between successive terms is

$$r_n = \frac{a_{n+1}}{a_n} = \frac{\lambda}{(n+1)^\nu}.$$

For $\nu > 0$, we have $r_n \to 0$ as $n \to \infty$, which guarantees the convergence of the series. Moreover, $(r_n)_{n \geq 0}$ is monotonically decreasing, ensuring that the conditions for the applicability of the Ratio-Bounding Pair method are satisfied.

- Sellers, 2021
- Conway & Maxwell, 1962
- Carvalho, Silva, and Moreira, 2025

Ratio Bounding Pair for COM-Poisson

FGV EMAp
ESCOLA DE MATEMÁTICA APLICADA

Algorithm: Ratio Bounding Pair

Initialize a_0, a_1 and $n = 1$
while $a_n > a_{n-1}$ **or** $a_n \left(1 - \frac{a_n}{a_{n-1}}\right)^{-1} \geq 2\varepsilon$ **do**
  $n \leftarrow n + 1$
  Evaluate a_n
end while
Evaluate $S_n = \sum_{k=0}^{n} a_k$
Return $S_n + \frac{a_n}{2}\left(\frac{L}{1-L} + \left(1 - \frac{a_n}{a_{n-1}}\right)^{-1}\right)$

- In code, we use log-scale and logsumexp to prevent numerical errors.
- The normalisation constant of the Poisson-COM distribution starts increasing and from the term $\lfloor \lambda^{1/\nu} \rfloor$ it is decreasing.

▶ Carvalho, Silva, and Moreira, 2025

Case Study 1: inventory dataset

The utility of the COM-Poisson model is illustrated using real-world count data, as an alternative to the Poisson model.

- **Dataset:** Quarterly sales counts of a specific clothing item (with overdispersion).
- **Model Comparison:** The COM-Poisson provides the best visual fit, especially for low counts, outperforming the Negative Binomial. While the Poisson captures the mean, it fails to account for overdispersion and, thus, underestimates variability.

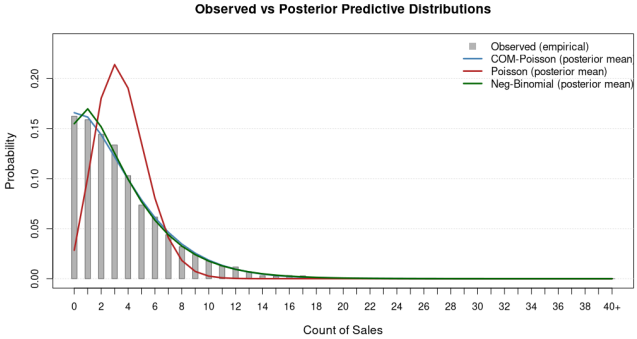

▶ Shmueli et al., 2005

Sensitivity to Numerical Precision

How does the precision level affect sampling algorithms, like the MCMC?

Sensitivity to Numerical Precision

FGV EMAp
ESCOLA DE MATEMÁTICA APLICADA

How does the precision level affect sampling algorithms, like the MCMC?

We investigate how the numerical precision used to approximate the COM-Poisson normalising constant affects MCMC performance.

- **Key Component:**
  - The COM-Poisson PMF depends on $Z_\varepsilon(\lambda, \nu)$, computed numerically with precision ε
- **Experimental Design:**
  - 100 independent MCMC runs for each ε
  - Each run:
    - 1,000 warm-up iterations
    - 500 sampling iterations
- **Metrics**
  - Effective Sample Size (ESS)
  - Convergence diagnostic $\hat{R}$
  - Execution time and ESS/sec

Case Study 1: inventory dataset

- **Stable Convergence:** For high-precision settings ($\varepsilon = 2^{-32}, 2^{-16}, 2^{-8}$), MCMC convergence is stable.
- **Best Balance:** $\varepsilon = 2^{-8}$ provides good sampling efficiency and the lowest execution time among the high-precision settings.
- **Collapse at Low Precision:** At $\varepsilon = 2^{-4}$, convergence deteriorates severely, indicating that adequate numerical precision is essential for reliable inference in this highly over-dispersed dataset.

ε	ESS Mean	$\hat{R}$ Mean	Mean Time (sec)	ESS/sec
2^{-32}	93.6	1.016	1210.9	0.077
2^{-16}	**94.0**	**1.014**	776.3	0.121
2^{-8}	84.1	1.018	**522.9**	**0.163**
2^{-4}	5.5	1.577	22666.2	0.000

Case Study 2: house prices dataset

- **Dataset**: bathroom counts in residential properties (with underdispersion)
- **Model Comparison:** The COM-Poisson model provides the closest visual agreement with the empirical distribution of the bathroom count. In contrast, both the Poisson and Neg-Binomial models, while reaching the average of the distribution, fail to accurately fit the specific shape of the data.

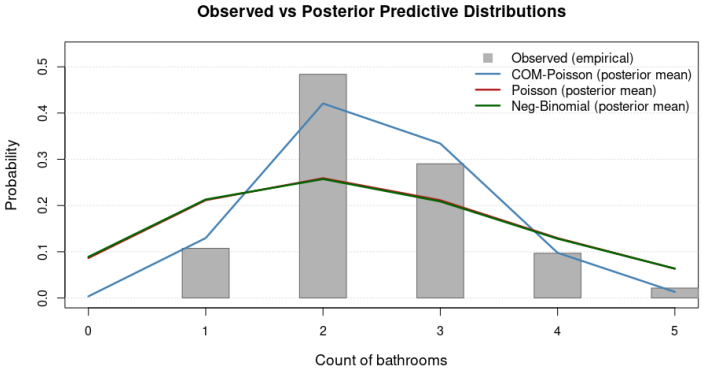

Case Study 2: house prices dataset

FGV EMAp
ESCOLA DE MATEMÁTICA APLICADA

- **Stable Convergence**: For high-precision settings ($\varepsilon = 2^{-32}, 2^{-16}, 2^{-8}, 2^{-4}$) , MCMC convergence is stable.
- **Best Balance**: $\varepsilon = 2^{-4}$ provides the highest sampling efficiency and lower execution time among the high-precision settings.
- **Collapse at Low Precision**: At $\varepsilon = 2^{-1}, 2^{-2}$, convergence deteriorates severely.

ε	ESS mean	Rhat mean	Time mean (sec)	ESS/sec
2^{-32}	330.4	1.007	41.1	**8.047**
2^{-16}	323.1	1.006	**41.0**	7.872
2^{-8}	327.7	1.006	42.3	7.752
2^{-4}	**345.6**	**1.005**	43.3	7.990
2^{-2}	308.0	1.055	65.0	4.735
2^{-1}	263.4	1.155	101.0	2.609

The Challenge of Model Correctness

The validation of statistical inference procedures is crucial. Theoretical models and their computational implementations can diverge, leading to misleading results.

- **Discrepancies:** Can arise from numerical approximations, implementation errors, or unverified assumptions.
- **Impact:** Biased estimates, incorrect uncertainty quantification, and flawed scientific conclusions.

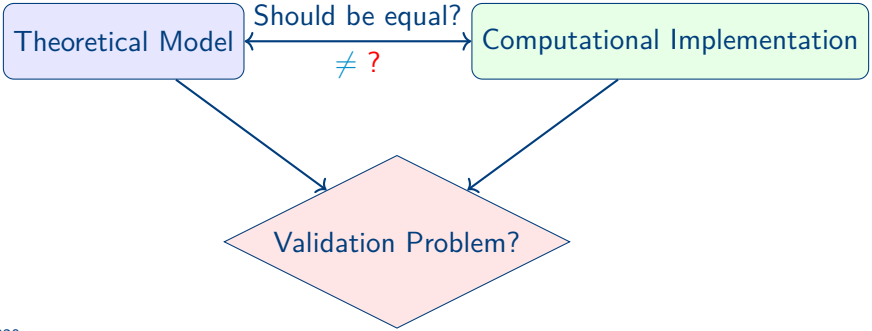

▶ Talts et al., 2020

SBC: A Framework for Validating Bayesian Inference

FGV EMAp
ESCOLA DE MATEMÁTICA APLICADA

Simulation-Based Calibration (SBC) is a robust methodology for validating Bayesian inference algorithms, ensuring that the posterior distributions are consistent with the generative model.

- **Central Principle:** If an MCMC sampler is correctly implemented and the model is valid, the distributions of the ranks of the true parameters within the posterior samples should be uniform.
- **Why it's Important:** It diagnoses systematic biases that traditional MCMC diagnostics (like $\hat{R}$ and ESS) may not detect.

- Talts et al., 2020
- Cook, Gelman, and Rubin, 2006

Simulation-Based Calibration (SBC) Flowchart

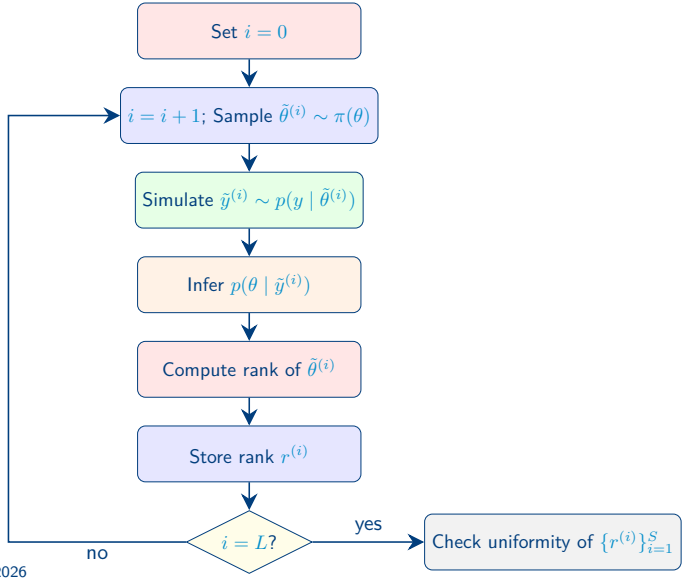

- Talts et al., 2020
- Michels's presentation, 2026

SBC Algorithm

Algorithm: Simulation-Based Calibration Procedure

Require: Inference Algorithm $\mathcal{M}$ with prior $\pi(\theta)$ and likelihood $\pi(y \mid \theta)$; Number of simulations N; Number of posterior samples L.

for $i = 1$ to N **do**

Sample true parameter: $\tilde{\theta}^{(i)} \sim \pi(\theta)$

Simulate data: $\tilde{y}^{(i)} \sim \pi\left(y \mid \tilde{\theta}^{(i)}\right)$

Run inference algorithm to obtain L posterior samples: $\theta_1^{(i)}, \cdots, \theta_L^{(i)} \sim \mathcal{M}(\tilde{y}^{(i)})$

Calculate the rank: $r^{(i)} = \sum_{l=1}^{L} \mathbb{I}\left(\theta_l^{(i)} < \tilde{\theta}^{(i)}\right)$

end for

return $\left\{r^{(i)}\right\}_{i=1}^{N}$

- Talts et al., 2020

Simulation-Based Calibration (SBC)

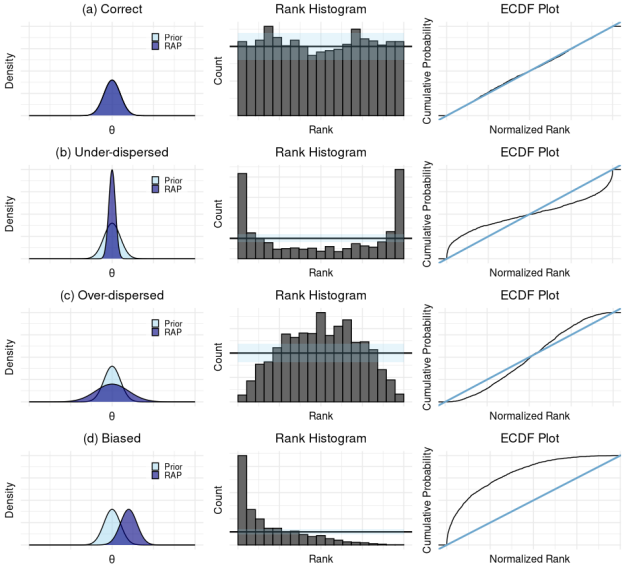

- Mendes et al., 2024

SBC for COM-Poisson Models: Requirements

To perform Simulation-Based Calibration effectively for COM-Poisson models, several key components are required to ensure the validity of the calibration process.

- **Sampler for Data Generation:** A reliable sampler for generating data from the COM-Poisson distribution ($\tilde{y} \sim \pi(y \mid \tilde{\theta})$).
- **Well-Defined Prior:** A suitable prior distribution ($\pi(\theta)$) for the COM-Poisson parameters (μ, ν) (with reparametrisation $\mu := \lambda^{1/\nu}$) must be specified.
- **Inference Algorithm:** The Bayesian inference algorithm to be tested.

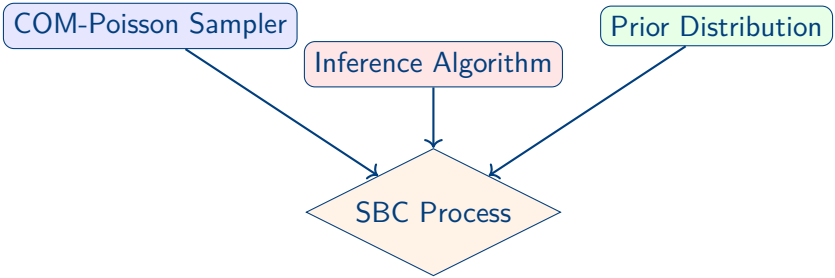

▶ Talts et al., 2020

Sampling from the COM-Poisson Distribution

Sampling from the COM-Poisson distribution is essential for various tasks, including data simulation and Monte Carlo-based inference.

- **Ratio-of-Uniform Method:** A general method that allows sampling from a distribution $\pi(x) = \tilde{\pi}(x)/Z$ by sampling uniform variables.
- **Rejection Sampling:** Methods tailored for the COM-Poisson, often using Poisson or geometric distributions as envelopes.
- **Exact Simulation with Ratio Bounding Pair:** This method, adapted from the Ratio Bounding Pair approach, is particularly important as it ensures no approximation error in the data generation phase of SBC.

▶ Benson and Friel, 2021

Sampling from COM-Poisson Distribution

FGV EMAp
ESCOLA DE MATEMÁTICA APLICADA

Algorithm: Exact Simulation with Ratio Bounding Pair

Simulate $z \sim U(0,1)$
Initialize $k = 0$, $j = 0$
Compute S_j and S_k
Compute (L_k, U_k)
Compute $(\underline{F}_{j,k}, \overline{F}_{j,k})$
while $\underline{F}_{j,k} < z$ **do**
  if $\underline{F}_{j,k} < z < \overline{F}_{j,k}$ **then**
    Update $k = k + 1$
    Compute $a[k] = \tilde{\pi}(k)$
    Update $S_k = \mathrm{sum}(a)$
    Compute (L_k, U_k)
  else
    Update $j = j + 1$
    Update $S_j = \mathrm{sum}(a[1:j])$
  end if
  Compute $(\underline{F}_{j,k}, \overline{F}_{j,k})$
end while
Return j

Choice of Prior for the COM-Poisson

Defining a suitable prior for the COM-Poisson parameters (μ, ν) (with reparametrisation $\mu := \lambda^{1/\nu}$) is crucial in Bayesian inference. In the context of Simulation-Based Calibration (SBC), the prior must be flexible and avoid overly informative assumptions.

- **In literature:** Conjugate priors have been proposed for the COM-Poisson model, but they are generally unsuitable for SBC. Existing works often suggest Gamma priors for both μ and ν.
- **Our choice:**
  - To balance under- and overdispersion: $\nu \sim \mathrm{Gamma}(1, 1)$
  - For flexibility across a wide range of applications:

$$\mu \sim \mathrm{Gamma}(\alpha, \beta), \quad \alpha, \beta > 0$$

  - In practice: $\mu \sim \mathrm{Gamma}(0.1, 0.01)$. With Mean $= 10$ and Variance $= 1000$.

- Mikkola et al, 2024
- Benson and Friel, 2021
- Boatwright et al., 2006
- Meyer, Graye, and Sellers (2023)

SBC Experimental Configuration

The Simulation-Based Calibration (SBC) experiments were meticulously designed to assess the computational faithfulness of the COM-Poisson inference algorithm across various precision levels.

- **Number of SBC Simulations (N):** 1,000 replications.
- **Number of Posterior Samples (L):** 200 posterior samples for rank calculation per replication.
- **Observations per Dataset:** 100 observations per simulated dataset.
- **Warmup Iterations:** 1,000 warmup iterations for the MCMC sampler.
- **Inference Algorithm:** MCMC implemented in Stan, utilising adaptive truncation for the COM-Poisson normalising constant.
- **Precision Levels (ε):** Six distinct levels of numerical precision were evaluated: $2^{-53}, 2^{-32}, 2^{-16}, 2^{-8}, 2^{-4}, 2^{-2}$.

▶ Talts et al., 2020

SBC Rank Histograms: Mean Parameter (μ)

The rank histograms for the mean parameter (μ) visually demonstrate how numerical precision affects calibration.

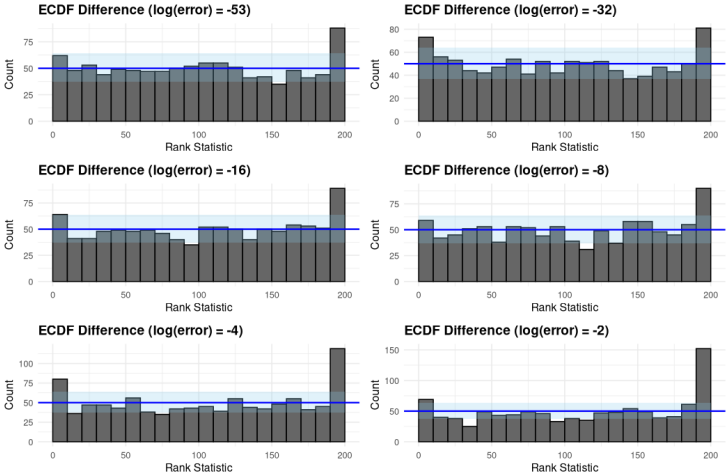

SBC Rank Histograms: Dispersion Parameter (ν)

Similarly, the rank histograms for the dispersion parameter (ν) confirm the impact of precision on calibration.

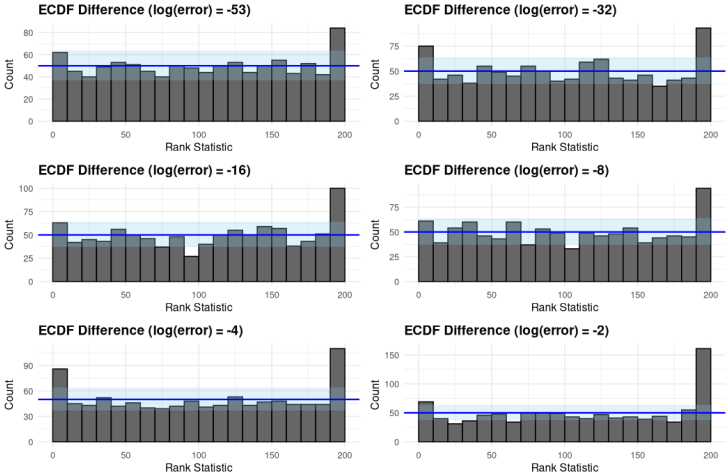

ECDF Diagnostics: Mean Parameter (μ)

The ECDF difference plots for μ provide a more granular view of the deviations from uniformity, complementing the rank histograms.

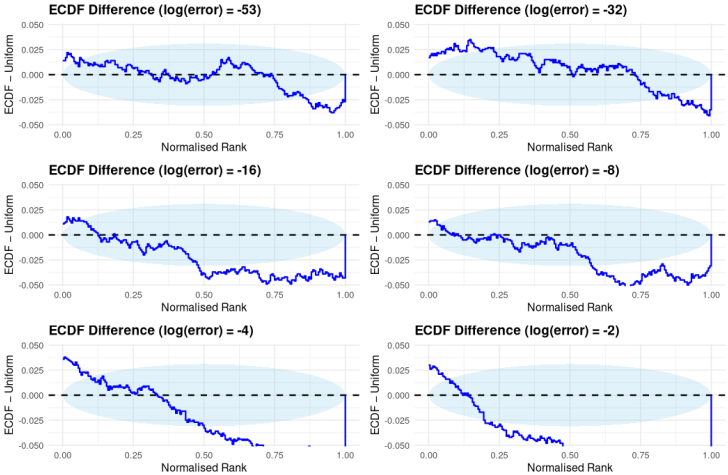

ECDF Diagnostics: Dispersion Parameter (ν)

The ECDF difference plots for ν reinforce the conclusions obtained for μ, highlighting the sensitivity of both parameters to numerical precision.

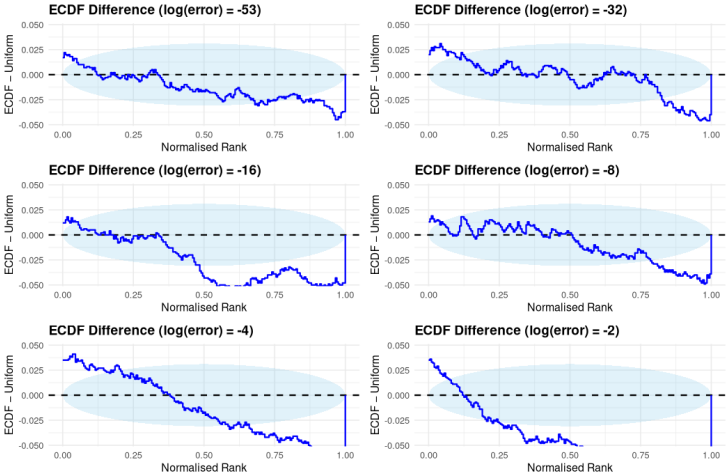

Future Work

- **Acceptance Probability Analysis**. Directly analysing how numerical errors in $Z(\lambda, \nu)$ propagate through the Metropolis-Hastings acceptance ratio, with the objective of understanding how these errors affect the behaviour of the MCMC algorithm.
- **Hybrid Approach**. Exploring hybrid methods combining asymptotic approximations with adaptive truncation.
- **Discrete Distribution Extensions**. Extending the adaptive truncation approach to a broader class of series could enable a wider range of applications.
- **Zero-Inflated Extensions**. Extending the computationally faithful COM-Poisson framework to include zero-inflated (ZICOMP) or hurdle variants to address datasets with excess zeros.
- **SBC Extensions**. Leveraging more advanced versions of SBC, such as Posterior SBC, to diagnose subtle forms of model misspecification or computational bias, particularly within hierarchical COM-Poisson models.

Thanks!

Any thoughts?

Special thanks to:

My supervisor (for patience!)
Coffee (for survival)
My family and friends (for everything)